

Urban Safety: Implementing Video Segmentation and Object Detection in Python and Building a Live Website

Group 1: Data Science

Bett Kipkemoi

Dominic Ong'aro

Faith Ngina

Dennis Mmenya Mwange

Video analytics pipeline for public safety

Features:

- Simulates simple synthetic public-space video with moving "people" and "vehicles".
- Runs motion segmentation (background subtraction) and object detection (Faster R-CNN if available, falls back to OpenCV HOG for people detection).
- Simple centroid tracker to detect loitering, crowding, and stationary vehicles (possible illegal parking).
- Produces an annotated output video, per-frame CSV log, and a short JSON report with actionable strategies.

Usage:

- Simulate and process:

```
python sentiment-analysis.py --simulate --out annotated.mp4 --log events.csv --report report.json
```

- Process an existing video:

```
python sentiment-analysis.py --input input.mp4 --out annotated.mp4 --log events.csv --report report.json
```

Dependencies:

- python >= 3.8
- opencv-python
- numpy
- pandas
- torch, torchvision

Note: This is a prototype for research/teaching. For production use, use optimized detectors, robust trackers, privacy-preserving pipelines, and follow local legal/regulatory guidance.

Reporting to the Website: Operational Logic

1. Data Acquisition

High-definition video feeds are ingested from a network of 1,200+ CCTV nodes distributed across Nairobi's Central Business District. Feeds are encrypted and transmitted via secure fiber optic links to the central processing unit.

2. Computer Vision Processing

Real-time AI models process video frames at 30fps.

- **Object Detection:** Identifying pedestrians, vehicles, and unattended objects.
- **Density Mapping:** Calculating crowd volume per square meter.
- **Anomaly Detection:** Flagging rapid movements, loitering, or traffic disruptions.\

3. Analysis and Correlation

Detected events are cross-referenced with historical data and temporal patterns. The system identifies deviations from the norm (e.g., a crowd forming at an unusual hour) to assess potential safety risks.

4. Response and Alerting

Verified incidents trigger automated alerts to the Command Center dashboard. High-priority threats (Tier-1) can be configured to automatically dispatch notifications to relevant field units or emergency services.

Daily Recommendations

Zone	Observation	Recommended Action	Priority
Moi Avenue	Crowd density > 85% capacity	Deploy crowd control unit; divert pedestrian flow	High
City Market	Increased petty theft reports (14:00-16:00)	Increase visible patrol presence during peak hours	Medium
Uhuru Highway	Traffic signal latency causing gridlock	Recalibrate signal timing software	Low

Site Overview: <https://urban-watch-ai-bettkipkemoi.replit.app>

SENTINELLE
NAIROBI SAFETY GRID

Command Center Live
Nairobi Central Business District • Zone A-1

System Status: Optimal | 3 Active Alerts

Active Cameras 1,248 (+12 new devices)

Daily Footfall 452k (+5% vs yesterday)

Safety Index 94.2 (-0.8% slight dip)

Incidents Today 23 (-5 below avg)

LIVE GEOSPATIAL VIEW
Heatmap Analysis: Moi Avenue
Updated 2m ago
Low Density | High Density

CAM-842 MOI AVENUE / KENYATTA AVE
Vehicle 91% Person 85% Group 75%
3 Objects Normal Activity 22:21:32

CAM-118 CITY MARKET ENTRANCE
Person 88% Group 95%
2 Objects Normal Activity 22:21:32

SENTINELLE
NAIROBI SAFETY GRID

Live Matrix
Real-time surveillance grid

Configure Grid

CAM-842 MOI AVENUE
Vehicle 98% Person 93% Group 87%
3 Objects Normal Activity 22:23:37

CAM-118 CITY MARKET
Person 96% Group 95%
2 Objects Normal Activity 22:23:37



CAM-284 UHURU HIGHWAY JCT
Person 92% Group 92%
2 Objects Normal Activity 22:23:37

CAM-889 WESTLANDS RD
Person 99% Group 97%
2 Objects Normal Activity 22:23:37

CAM-186 KIBERA ENTRY
Vehicle 93% Person 91% Group 88%
3 Objects Normal Activity 22:23:37

CAM-332 CBD CENTRAL
Person 95% Group 86%
2 Objects Normal Activity 22:23:37


SYSTEM ONLINE
Latency: 12ms




reports.json

```
{  
  "video": "simulated_input.mp4",  
  "processed_frames": 50,  
  "total_people_observations": 0,  
  "total_vehicle_observations": 0,  
  "hotspots": [],  
  "recommendations": [  
    "Increase patrol presence and lighting in identified hotspots (grid cells with high people counts).",  
    "Install time-limited parking restrictions and signage at bus_stop / market areas; enforce to reduce illegal parking.",  
    "Deploy crowd management protocols (additional personnel, temporary barriers) during peak times identified by spikes in crowd alerts.",  
    "Use real-time alerts to dispatch response teams for loitering and stationary vehicles, and integrate with local traffic enforcement.",  
    "Complement camera analytics with anonymization (blurring faces) and data governance policies to protect privacy."  
  ]  
}
```

annotated.mp4



simulated.mp4



Code:

```
import argparse
import cv2
import numpy as np
import time
import os
```

```
import csv
import json
from collections import deque, defaultdict
import torch
import torchvision
from torchvision import transforms
import pandas as pd

"""
Video analytics pipeline for public safety
"""

try:
    TORCH_AVAILABLE = True
except Exception:
    TORCH_AVAILABLE = False

# -----
# Simulation
# -----
def simulate_video(path="simulated.mp4", n_frames=600, size=(640, 480), fps=20):
    """Generate a synthetic public-space video with moving rectangles."""
    w, h = size
    fourcc = cv2.VideoWriter_fourcc(*"mp4v")
    out = cv2.VideoWriter(path, fourcc, fps, (w, h))

    rng = np.random.RandomState(42)
    people = []
    vehicles = []

    # spawn some initial actors
    for _ in range(8):
        people.append({
            "pos": np.array([rng.randint(0, w), rng.randint(0, h)]).astype(float),
            "vel": (rng.rand(2) - 0.5) * 4,
            "size": rng.randint(8, 16),
            "color": (255, 200, 200)
```

```

        })

for _ in range(3):
    vehicles.append({
        "pos": np.array([rng.randint(0, w), rng.randint(0, h)]).astype(float),
        "vel": (rng.rand(2) - 0.5) * 6,
        "size": rng.randint(24, 40),
        "color": (200, 255, 200)
    })

# define some zones (x1,y1,x2,y2) for "no-parking" or bus stops
zones = {
    "market": (int(0.1*w), int(0.1*h), int(0.35*w), int(0.35*h)),
    "bus_stop": (int(0.6*w), int(0.7*h), int(0.95*w), int(0.95*h)),
}

for f in range(n_frames):
    frame = np.full((h, w, 3), 30, dtype=np.uint8)

    # draw zones lightly
    for name, (x1,y1,x2,y2) in zones.items():
        cv2.rectangle(frame, (x1,y1), (x2,y2), (60,60,80), 1)

    # spawn occasional new people
    if rng.rand() < 0.03:
        people.append({
            "pos": np.array([rng.randint(0, w), rng.randint(0, h)]).astype(float),
            "vel": (rng.rand(2) - 0.5) * 4,
            "size": rng.randint(8, 16),
            "color": (255, 200, 200)
        })

    # update and draw people
    for p in people:
        p["pos"] += p["vel"]

        # bounce at borders
        for i in (0,1):
            if p["pos"][i] < 0 or p["pos"][i] >= (w if i==0 else h):
                p["vel"][i] *= -1
                p["pos"][i] = np.clip(p["pos"][i], 0, (w-1) if i==0 else (h-1))

        cx, cy = int(p["pos"][0]), int(p["pos"][1])

```

```

r = int(p["size"]/2)
cv2.circle(frame, (cx, cy), r, p["color"], -1)

# update and draw vehicles
for v in vehicles:
    # occasional stop in bus stop zone to simulate parking
    if np.random.rand() < 0.002:
        v["vel"] = np.array([0.0, 0.0])
    # random resume
    if np.random.rand() < 0.01 and np.linalg.norm(v["vel"]) < 0.5:
        v["vel"] = (rng.rand(2) - 0.5) * 6
    v["pos"] += v["vel"]
    for i in (0,1):
        limit = w if i==0 else h
        if v["pos"][i] < 0 or v["pos"][i] >= limit:
            v["vel"][i] *= -1
            v["pos"][i] = np.clip(v["pos"][i], 0, limit-1)
    cx, cy = int(v["pos"][0]), int(v["pos"][1])
    s = int(v["size"])
    cv2.rectangle(frame, (cx-s, cy-s//2), (cx+s, cy+s//2), v["color"], -1)

    # add synthetic background motion (e.g., trees shimmer)
    noise = (rng.rand(h, w) * 2).astype(np.uint8)
    frame = cv2.add(frame, np.stack([noise]*3, axis=2))

out.write(frame)

out.release()
return path

# -----
# Detector helpers
# -----
COCO_PERSON_ID = 1
COCO_VEHICLE_IDS = {3, 6, 8} # car, bus, truck (approx)
COCO_LABELS = None

def load_coco_labels():
    # short COCO label list (index starts at 1)
    labels = {1: "person", 2:"bicycle",3:"car",4:"motorcycle",5:"airplane",6:"bus",

```

```

    7:"train",8:"truck",9:"boat",10:"traffic light"})

return labels

def load_torch_detector(device="cpu"):
    global COCO_LABELS
    COCO_LABELS = load_coco_labels()
    if not TORCH_AVAILABLE:
        return None, device
    model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)
    model.to(device)
    model.eval()
    transform = transforms.Compose([
        transforms.ToTensor()])
    return (model, transform), device

# -----
# Simple centroid tracker
# -----
class CentroidTracker:

    def __init__(self, max_distance=50, max_lost=30):
        self.next_id = 1
        self.objects = dict() # id -> (centroid, last_seen_frame, trace deque)
        self.max_distance = max_distance
        self.max_lost = max_lost

    def update(self, detections, frame_idx):
        # detections: list of (x, y) centroids
        assigned = set()
        new_objects = dict()

        # simple greedy assignment
        for oid, (centroid, last_seen, trace) in list(self.objects.items()):
            best = None
            best_d = None
            for i, c in enumerate(detections):
                if i in assigned:
                    continue
                d = np.linalg.norm(np.array(centroid) - np.array(c))
                if d <= self.max_distance and (best is None or d < best_d):

```

```

        best = i
        best_d = d

    if best is not None:
        c = detections[best]
        trace.append(c)
        if len(trace) > 64:
            trace.popleft()
        new_objects[oid] = (c, frame_idx, trace)
        assigned.add(best)

    else:
        # not matched: keep but mark lost
        if frame_idx - last_seen <= self.max_lost:
            new_objects[oid] = (centroid, last_seen, trace)
        # else drop

# add unmatched detections
for i, c in enumerate(detections):
    if i in assigned:
        continue
    oid = self.next_id
    self.next_id += 1
    dq = deque([c], maxlen=64)
    new_objects[oid] = (c, frame_idx, dq)

self.objects = new_objects
return self.objects

# -----
# Processing pipeline
# -----


def process_video(input_path, output_path="annotated.mp4", log_path="events.csv", report_path="report.json",
                  device="cpu", run_detector=True, detect_every=1):

    # try to load detector
    detector, device = None, device

    if run_detector:
        det_res = load_torch_detector(device=device)
        if det_res[0] is None:
            detector = None
            print("Torch unavailable: falling back to HOG people detector.")
        else:

```

```

        detector = det_res[0]
    else:
        detector = None

    cap = cv2.VideoCapture(input_path)
    if not cap.isOpened():
        raise FileNotFoundError(f"Cannot open video {input_path}")

    fps = cap.get(cv2.CAP_PROP_FPS) or 20.0
    w = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
    h = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))
    fourcc = cv2.VideoWriter_fourcc(*"mp4v")
    out = cv2.VideoWriter(output_path, fourcc, fps, (w,h))

    # background subtractor
    backSub = cv2.createBackgroundSubtractorMOG2(history=500, detectShadows=True)

    # tracker for people and vehicles (separate)
    person_tracker = CentroidTracker(max_distance=60, max_lost=int(fps*3))
    vehicle_tracker = CentroidTracker(max_distance=80, max_lost=int(fps*10))

    # prepare HOG fallback
    hog = None
    if detector is None:
        hog = cv2.HOGDescriptor()
        hog.setSVMClassifier(cv2.HOGDescriptor_getDefaultPeopleDetector())

    frame_idx = 0
    logs = []
    hotspots = defaultdict(int)

    while True:
        ret, frame = cap.read()
        if not ret:
            break
        frame_idx += 1
        t0 = time.time()

        fgmask = backSub.apply(frame)
        # morphological cleanup

```

```

fg = cv2.medianBlur(fgmask, 5)
_, fg = cv2.threshold(fg, 200, 255, cv2.THRESH_BINARY)
fg = cv2.morphologyEx(fg, cv2.MORPH_CLOSE, np.ones((5,5), np.uint8))

people_centroids = []
vehicle_centroids = []
detections_out = []

if detector and frame_idx % detect_every == 0:
    model, transform = detector
    img_t = transform(frame).to(device)
    with torch.no_grad():
        preds = model([img_t])[0]
        boxes = preds["boxes"].cpu().numpy()
        labels = preds["labels"].cpu().numpy()
        scores = preds["scores"].cpu().numpy()

    for box, label, score in zip(boxes, labels, scores):
        if score < 0.6:
            continue
        x1,y1,x2,y2 = box.astype(int)
        cx = int((x1+x2)/2); cy = int((y1+y2)/2)
        lbl = COCO_LABELS.get(int(label), str(int(label)))
        detections_out.append((lbl, score, (x1,y1,x2,y2)))
        if int(label) == COCO_PERSON_ID:
            people_centroids.append((cx, cy))
        if int(label) in COCO_VEHICLE_IDS:
            vehicle_centroids.append((cx, cy))
    else:
        # fallback: use background segments and simple contours; detect people via HOG
        contours, _ = cv2.findContours(fg, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
        for cnt in contours:
            x,y,ww,hh = cv2.boundingRect(cnt)
            area = ww*hh
            if area < 200: continue
            cx = x+ww//2; cy = y+hh//2
            # classify by size heuristics
            if hh < 30:
                people_centroids.append((cx, cy))
                detections_out.append(("person", 0.5, (x,y,x+ww,y+hh)))

```

```

        else:
            vehicle_centroids.append((cx, cy))
            detections_out.append(("vehicle", 0.5, (x,y,x+ww,y+hh)))

# HOG refine people (optional)
if hog is not None:
    rects, weights = hog.detectMultiScale(frame, winStride=(8,8), padding=(8,8), scale=1.05)
    for (x,y,ww,hh), wgt in zip(rects, weights):
        if wgt < 0.6: continue
        cx = int(x+ww/2); cy = int(y+hh/2)
        people_centroids.append((cx, cy))
        detections_out.append(("person_hog", float(wgt), (x,y,x+ww,y+hh)))

# update trackers
persons = person_tracker.update(people_centroids, frame_idx)
vehicles = vehicle_tracker.update(vehicle_centroids, frame_idx)

alerts = []

# detect loitering: person whose trace is within small radius for long time
for pid, (centroid, last_seen, trace) in persons.items():
    if len(trace) >= int(fps*10): # been tracked for >10s
        pts = np.array(trace)
        dmax = np.max(np.linalg.norm(pts - pts[0], axis=1))
        if dmax < 20:
            alerts.append({"type": "loitering", "id": pid, "pos": centroid})
            # annotate
            cv2.putText(frame, f"LOITER-{pid}", (centroid[0]+5, centroid[1]-5),
                       cv2.FONT_HERSHEY_SIMPLEX, 0.5, (0,0,255), 2)

# stationary vehicle detection -> possible illegal parking
for vid, (centroid, last_seen, trace) in vehicles.items():
    if len(trace) >= int(fps*8):
        pts = np.array(trace)
        dmax = np.max(np.linalg.norm(pts - pts[0], axis=1))
        if dmax < 10:
            alerts.append({"type": "illegal_parking", "id": vid, "pos": centroid})
            cv2.putText(frame, f"PARK-{vid}", (centroid[0]+5, centroid[1]-5),
                       cv2.FONT_HERSHEY_SIMPLEX, 0.5, (0,165,255), 2)

```

```

# crowd detection: many people in small area
if len(people_centroids) >= 6:
    # compute pairwise density
    pts = np.array(people_centroids)
    if pts.shape[0] >= 6:
        # compute bounding box area
        x_min, y_min = pts.min(axis=0)
        x_max, y_max = pts.max(axis=0)
        area = (x_max - x_min + 1) * (y_max - y_min + 1)
        density = len(people_centroids) / max(area, 1)
        if density > 0.002: # heuristic density threshold
            alerts.append({"type": "crowd", "count": len(people_centroids), "bbox": [int(x_min), int(y_min), int(x_max), int(y_max)]})
            cv2.rectangle(frame, (int(x_min), int(y_min)), (int(x_max), int(y_max)), (0,0,255), 2)
            cv2.putText(frame, f"CROWD {len(people_centroids)}", (int(x_min), int(y_min)-8),
                       cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0,0,255), 2)

# annotate detections
for lbl, score, box in detections_out:
    x1,y1,x2,y2 = box
    color = (200,200,255) if "person" in lbl else (200,255,200)
    cv2.rectangle(frame, (x1,y1), (x2,y2), color, 1)
    cv2.putText(frame, f"[{lbl}]:{score:.2f}", (x1, max(0,y1-6)),
               cv2.FONT_HERSHEY_SIMPLEX, 0.4, color, 1)

# annotate trackers
for pid, (centroid, last_seen, trace) in persons.items():
    cv2.circle(frame, (int(centroid[0]), int(centroid[1])), 3, (255,0,0), -1)
    cv2.putText(frame, f"P{pid}", (int(centroid[0])+4, int(centroid[1])+4),
               cv2.FONT_HERSHEY_SIMPLEX, 0.4, (255,0,0), 1)
for vid, (centroid, last_seen, trace) in vehicles.items():
    cv2.circle(frame, (int(centroid[0]), int(centroid[1])), 3, (0,255,0), -1)
    cv2.putText(frame, f"V{vid}", (int(centroid[0])+4, int(centroid[1])+4),
               cv2.FONT_HERSHEY_SIMPLEX, 0.4, (0,255,0), 1)

# write frame
out.write(frame)

# logging
ts = frame_idx / fps

```

```

entry = {
    "frame": frame_idx,
    "time_s": round(ts,2),
    "people_count": len(people_centroids),
    "vehicle_count": len(vehicle_centroids),
    "alerts": alerts
}
logs.append(entry)

# hotspot accumulation by simple grid
for c in people_centroids:
    gx = int(c[0] / (w/10))
    gy = int(c[1] / (h/10))
    hotspots[(gx,gy)] += 1

if frame_idx % int(fps*10) == 0:
    print(f"[{frame_idx}] people:{len(people_centroids)} vehicles:{len(vehicle_centroids)}"
          f"\nalerts:{len(alerts)} dt:{(time.time()-t0):.2f}s")

cap.release()
out.release()

# save logs
df = pd.DataFrame([
    {"frame": e["frame"], "time_s": e["time_s"], "people_count": e["people_count"],
     "vehicle_count": e["vehicle_count"], "alerts": json.dumps(e["alerts"])}
    for e in logs
])
df.to_csv(log_path, index=False)

# summarize hotspots top-k
hot = sorted(hotspots.items(), key=lambda x: -x[1])[:8]
hot_zones = []
for (gx,gy), cnt in hot:
    hot_zones.append({"grid": [int(gx),int(gy)], "count": int(cnt)})

# generate simple report with actionable strategies
report = {
    "video": os.path.basename(input_path),
    "processed_frames": frame_idx,
}

```

```

        "total_people_observations": int(sum([e["people_count"] for e in logs])),
        "total_vehicle_observations": int(sum([e["vehicle_count"] for e in logs])),
        "hotspots": hot_zones,
        "recommendations": [
            "Increase patrol presence and lighting in identified hotspots (grid cells with high people counts).",
            "Install time-limited parking restrictions and signage at bus_stop / market areas; enforce to reduce illegal parking.",
            "Deploy crowd management protocols (additional personnel, temporary barriers) during peak times identified by spikes in crowd alerts.",
            "Use real-time alerts to dispatch response teams for loitering and stationary vehicles, and integrate with local traffic enforcement.",
            "Complement camera analytics with anonymization (blurring faces) and data governance policies to protect privacy."
        ]
    }

    with open(report_path, "w") as f:
        json.dump(report, f, indent=2)

    print(f"Processing complete. Annotated video: {output_path}, logs: {log_path}, report: {report_path}")
    return output_path, log_path, report_path

# -----
# CLI
# -----
def main():

    p = argparse.ArgumentParser(description="Video analytics prototype for urban public safety.")
    p.add_argument("--simulate", action="store_true", help="Create a simulated video and process it.")
    p.add_argument("--input", type=str, help="Input video path.")
    p.add_argument("--out", type=str, default="annotated.mp4", help="Annotated output video path.")
    p.add_argument("--log", type=str, default="events.csv", help="CSV log output.")
    p.add_argument("--report", type=str, default="report.json", help="JSON report output.")
    p.add_argument("--device", type=str, default="cpu", help="torch device (cpu or cuda).")
    p.add_argument("--frames", type=int, default=600, help="Frames to simulate when --simulate.")

    args = p.parse_args()

    if args.simulate:
        sim_path = "simulated_input.mp4"
        print("Generating simulated video:", sim_path)
        simulate_video(path=sim_path, n_frames=args.frames)

```

```
    input_path = sim_path

else:
    if not args.input:
        p.print_help()
        return

    input_path = args.input

# try to use GPU if requested
device = args.device
if device != "cpu" and TORCH_AVAILABLE and torch.cuda.is_available():
    device = "cuda"
else:
    device = "cpu"

process_video(input_path, output_path=args.out, log_path=args.log, report_path=args.report,
device=device)

if __name__ == "__main__":
    main()
```