

mini-project

January 4, 2026

1 Decision Tree Algorithms and Ensemble Methods

In this mini-project, Decision Tree algorithms and Ensemble Methods:- 1. Bagging (random forest)
2. Boosting (ada boost), and 3. Stacking (stacking classifier)

are applied to the Credit Card Fraud Detection dataset [<https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>] and the performance of each method compared. The dataset contains credit card transactions, where the focus is to predict whether a transaction is fraudulent or not.

The steps undertaken are: - Step 1: Data Exploration and Visualization - Step 2: Data Preprocessing - Step 3: Implementation of Decision Tree Classifier - Step 4: Implementation of Random Forest Classifier (Bagging) - Step 5: Implementation of AdaBoost Classifier (Boosting) - Step 6: Implementation of Stacking Classifier (Stacking) - Step 7: Model Evaluation and Comparison

1.1 Step 1: Data Exploration and Visualization

Libraries for data manipulation, model training and visualization are imported.

```
[1]: # import libraries
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```



```
[2]: # load the dataset
df = pd.read_csv('/Users/bett/downloads/creditcard.csv')
df.head(5)
```


	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	

```

1 0.085102 -0.255425 ... -0.225775 -0.638672 0.101288 -0.339846 0.167170
2 0.247676 -1.514654 ... 0.247998 0.771679 0.909412 -0.689281 -0.327642
3 0.377436 -1.387024 ... -0.108300 0.005274 -0.190321 -1.175575 0.647376
4 -0.270533 0.817739 ... -0.009431 0.798278 -0.137458 0.141267 -0.206010

```

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

1.1.1 Explore the Dataset

- displaying basic information
- analyzing the distribution of target variable (fraudulent or not)
- visualizing the relationships between features using correlation matrices, pairplots and histograms

```
[3]: # information of the variables
print("Information of the variables")
print(df.info())
```

```
Information of the variables
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   Time      284807 non-null   float64
 1   V1        284807 non-null   float64
 2   V2        284807 non-null   float64
 3   V3        284807 non-null   float64
 4   V4        284807 non-null   float64
 5   V5        284807 non-null   float64
 6   V6        284807 non-null   float64
 7   V7        284807 non-null   float64
 8   V8        284807 non-null   float64
 9   V9        284807 non-null   float64
 10  V10       284807 non-null   float64
 11  V11       284807 non-null   float64
 12  V12       284807 non-null   float64
 13  V13       284807 non-null   float64
 14  V14       284807 non-null   float64
 15  V15       284807 non-null   float64
 16  V16       284807 non-null   float64
```

```

17 V17      284807 non-null  float64
18 V18      284807 non-null  float64
19 V19      284807 non-null  float64
20 V20      284807 non-null  float64
21 V21      284807 non-null  float64
22 V22      284807 non-null  float64
23 V23      284807 non-null  float64
24 V24      284807 non-null  float64
25 V25      284807 non-null  float64
26 V26      284807 non-null  float64
27 V27      284807 non-null  float64
28 V28      284807 non-null  float64
29 Amount   284807 non-null  float64
30 Class    284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
None

```

```
[4]: # print basic statistics
print("Basic statistics of the dataset")
print(df.describe())
```

Basic statistics of the dataset

	Time	V1	V2	V3	V4	\
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	94813.859575	1.175161e-15	3.384974e-16	-1.379537e-15	2.094852e-15	
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	
	V5	V6	V7	V8	V9	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	1.021879e-15	1.494498e-15	-5.620335e-16	1.149614e-16	-2.414189e-15	
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00	
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01	
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01	
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02	
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01	
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01	
	V21	V22	V23	V24	\	
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	...	1.628620e-16	-3.576577e-16	2.618565e-16	4.473914e-15	
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01	
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00	

```

25% ... -2.283949e-01 -5.423504e-01 -1.618463e-01 -3.545861e-01
50% ... -2.945017e-02 6.781943e-03 -1.119293e-02 4.097606e-02
75% ... 1.863772e-01 5.285536e-01 1.476421e-01 4.395266e-01
max ... 2.720284e+01 1.050309e+01 2.252841e+01 4.584549e+00

V25 V26 V27 V28 Amount \
count 2.848070e+05 2.848070e+05 2.848070e+05 2.848070e+05 284807.000000
mean 5.109395e-16 1.686100e-15 -3.661401e-16 -1.227452e-16 88.349619
std 5.212781e-01 4.822270e-01 4.036325e-01 3.300833e-01 250.120109
min -1.029540e+01 -2.604551e+00 -2.256568e+01 -1.543008e+01 0.000000
25% -3.171451e-01 -3.269839e-01 -7.083953e-02 -5.295979e-02 5.600000
50% 1.659350e-02 -5.213911e-02 1.342146e-03 1.124383e-02 22.000000
75% 3.507156e-01 2.409522e-01 9.104512e-02 7.827995e-02 77.165000
max 7.519589e+00 3.517346e+00 3.161220e+01 3.384781e+01 25691.160000

Class
count 284807.000000
mean 0.001727
std 0.041527
min 0.000000
25% 0.000000
50% 0.000000
75% 0.000000
max 1.000000

```

[8 rows x 31 columns]

```
[5]: # print number of rows and columns
print("Number of rows and columns in the dataset:")
print(df.shape)
```

Number of rows and columns in the dataset:
(284807, 31)

```
[6]: # check for missing values
print("Missing values in the dataset:")
print(df.isnull().sum())
```

Missing values in the dataset:

Time	0
V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0

```
V10      0
V11      0
V12      0
V13      0
V14      0
V15      0
V16      0
V17      0
V18      0
V19      0
V20      0
V21      0
V22      0
V23      0
V24      0
V25      0
V26      0
V27      0
V28      0
Amount    0
Class     0
dtype: int64
```

```
[7]: # check for duplicate rows and remove them
duplicate_rows = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicate_rows}")
if duplicate_rows > 0:
    df = df.drop_duplicates()
    print(f"Duplicate rows removed. New shape of the dataset: {df.shape}")
```


```
Number of duplicate rows: 1081
Duplicate rows removed. New shape of the dataset: (283726, 31)
```

```
[8]: # analyze the target variable
print("Distribution of the target variable:")
print(df['Class'].value_counts())
```


```
Distribution of the target variable:
Class
0    283253
1      473
Name: count, dtype: int64
```

```
[9]: # visualize the relationship between features using correlation matrix
plt.figure(figsize=(12, 10))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Matrix')
```

```
plt.show()
```



```
[10]: # histograms of features V1 to V28
df.hist(column=[f'V{i}' for i in range(1, 29)], figsize=(20, 15))
plt.suptitle('Histograms of Features V1 to V28')
plt.show()
```



1.2 Step 2: Data Preprocessing

- Data Splitting
- Handling missing values
- Data Scaling

```
[11]: # split the dataset into training and testing sets
"""
Split the dataset into features and target variable first, then into training and testing sets.
"""
X = df.drop('Class', axis=1)
y = df['Class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

1.3 Step 3: Implementing Decision Tree Classifier

```
[37]: # handle missing values if any (not needed here as there are no missing values)
# data scaling (not needed here as Decision Trees are not affected by feature
    ↪scaling)

# train the Decision Tree model
model = DecisionTreeClassifier(random_state=42)
model.fit(X_train, y_train)

# make predictions
y_pred = model.predict(X_test)


# evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
#classification report
from sklearn.metrics import classification_report
dt_classification_report = classification_report(y_test, y_pred)
print("Decision Tree Classification Report:")
print(dt_classification_report)

print(f'Accuracy: {accuracy:.4f}')
# visualize the confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', u
    ↪xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion Matrix')
plt.show()
```

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56656
1	0.67	0.71	0.69	90
accuracy			1.00	56746
macro avg	0.83	0.86	0.84	56746
weighted avg	1.00	1.00	1.00	56746

Accuracy: 0.9990



1.4 Step 4: Implementing Random Forest Classifier (Bagging)

```
[35]: # data scaling for random forest
# import StandardScaler, RandomForestClassifier and classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# train the Random Forest model
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train_scaled, y_train)

# make predictions with Random Forest
y_rf_pred = rf_model.predict(X_test_scaled)

# evaluate the Random Forest model
rf_accuracy = accuracy_score(y_test, y_rf_pred)
rf_conf_matrix = confusion_matrix(y_test, y_rf_pred)
rf_classification_report = classification_report(y_test, y_rf_pred)
```

```

# print Random Forest results
print(f'Random Forest Accuracy: {rf_accuracy:.4f}')
print('Classification Report:')
print(rf_classification_report)


# visualize the Random Forest confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(rf_conf_matrix, annot=True, fmt='d', cmap='Greens', ▾
            xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Random Forest Confusion Matrix')
plt.show()

```

Random Forest Accuracy: 0.9995

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56656
1	0.97	0.73	0.84	90
accuracy			1.00	56746
macro avg	0.99	0.87	0.92	56746
weighted avg	1.00	1.00	1.00	56746



1.5 Step 5: Implementing AdaBoost Classifier (Boosting)

```
[14]: # use the scaled data
# train AdaBoost Classifier
from sklearn.ensemble import AdaBoostClassifier
adb_model = AdaBoostClassifier(random_state=42)
adb_model.fit(X_train_scaled, y_train)
```

```
[14]: AdaBoostClassifier(random_state=42)
```

```
[15]: # make predictions with AdaBoost
y_adb_pred = adb_model.predict(X_test_scaled)
```

```
[38]: # evaluate the AdaBoost model
adb_accuracy = accuracy_score(y_test, y_adb_pred)
adb_conf_matrix = confusion_matrix(y_test, y_adb_pred)
adb_classification_report = classification_report(y_test, y_adb_pred)
# print AdaBoost classification report
print("AdaBoost Classification Report:")
print(adb_classification_report)
```


AdaBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56656
1	0.79	0.68	0.73	90
accuracy			1.00	56746
macro avg	0.90	0.84	0.87	56746
weighted avg	1.00	1.00	1.00	56746

```
[33]: # print AdaBoost results
print(f'AdaBoost Accuracy: {adb_accuracy:.4f}')

# visualize the AdaBoost confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(adb_conf_matrix, annot=True, fmt='d', cmap='Greens',
            xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('AdaBoost Confusion Matrix')
plt.show()
```

AdaBoost Accuracy: 0.9992



1.6 Step 6: Implementing Stacking Classifier (Stacking)

```
[18]: #load packages
import pandas as pd
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import StackingRegressor
from sklearn.linear_model import LinearRegression
```

```
[19]: # initialize base learners
rf_regressor = RandomForestRegressor(n_estimators=100, random_state=42)
gb_regressor = GradientBoostingRegressor(n_estimators=100, random_state=42)

# create base learners
estimators = [
    ('rf', rf_regressor),
    ('gb', gb_regressor),
]
```

```
[20]: # create the stacking regressor using linear regression as the final estimator
stacking_model = StackingRegressor(
```

```
        estimators=estimators,
        final_estimator=LinearRegression()
)
```

```
[21]: # create the stacking model
stacking_model.fit(X_train_scaled, y_train)
#make predictions
y_pred_stacking = stacking_model.predict(X_test_scaled)
```

```
[22]: #evaluate the model
mse_stacking = mean_squared_error(y_test, y_pred_stacking)
r2 = stacking_model.score(X_test_scaled, y_test)
print(f'Stacking Regressor MSE: {mse_stacking:.4f}')
print(f'Stacking Regressor R^2: {r2:.4f}')
```

Stacking Regressor MSE: 0.0005
Stacking Regressor R²: 0.7078

1.7 Step 7: Compare All the Models

```
[43]: # compare all the models' performance
model_performance = {
    'Decision Tree Accuracy': accuracy,
    'Random Forest Accuracy': rf_accuracy,
    'AdaBoost Accuracy': adb_accuracy,
    'Stacking Regressor R^2': r2,    # Using R^2 score for stacking
    'Stacking Regressor MSE': mse_stacking
}

# print model performance in table format
for model, performance in model_performance.items():
    print(f'{model}: {performance:.4f}')

#print classification reports for all models
print("Decision Tree Classification Report:")
print(dt_classification_report)
print("Random Forest Classification Report:")
print(rf_classification_report)
print("AdaBoost Classification Report:")
print(adb_classification_report)
```

Decision Tree Accuracy: 0.9990
Random Forest Accuracy: 0.9995
AdaBoost Accuracy: 0.9992
Stacking Regressor R²: 0.7078
Stacking Regressor MSE: 0.0005
Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56656
1	0.67	0.71	0.69	90
accuracy			1.00	56746
macro avg	0.83	0.86	0.84	56746
weighted avg	1.00	1.00	1.00	56746

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56656
1	0.97	0.73	0.84	90
accuracy			1.00	56746
macro avg	0.99	0.87	0.92	56746
weighted avg	1.00	1.00	1.00	56746

AdaBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56656
1	0.79	0.68	0.73	90
accuracy			1.00	56746
macro avg	0.90	0.84	0.87	56746
weighted avg	1.00	1.00	1.00	56746

The report presents performance metrics for several machine learning models on a highly imbalanced binary classification dataset with 56,746 total samples, where the majority class (label 0) comprises 56,656 instances and the minority class (label 1) only 90 instances. All three classification models—Decision Tree, Random Forest, and AdaBoost—achieve near-perfect overall accuracy (0.9990, 0.9995, and 0.9992 respectively), but this is largely driven by excellent performance on the majority class, with precision, recall, and F1-score of 1.00 for class 0 across all models. For the critical minority class, performance varies significantly: Random Forest performs best with a precision of 0.97, recall of 0.73, and F1-score of 0.84; Decision Tree follows with precision 0.67, recall 0.71, and F1-score 0.69; while AdaBoost lags with precision 0.79, recall 0.68, and F1-score 0.73.

The inclusion of regression metrics—a Stacking Regressor with an R^2 of 0.7078 and a very low MSE of 0.0005—suggests that while the classification of categories is highly successful, the model also maintains a decent level of predictive power for continuous target variables, though there is still room for improvement in explaining the total variance.