

## Convolutional Neural Networks (CNNs)

This notebook demonstrates how to build a Convolutional Neural Network (CNN) to classify images from the Cifar10 dataset. We then evaluate the model on test data and report the performance metrics (accuracy, precision, recall).

```
#load necessary libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import itertools

import keras
from keras.datasets import cifar10
from tensorflow.keras import layers
from keras.models import Sequential
from sklearn.metrics import confusion_matrix, classification_report

%matplotlib inline

# load and preprocess data
(train_images, train_labels), (test_images, test_labels) = cifar10.load_data()

# normalize pixel values
train_images, test_images = train_images/255.0, test_images/255.0

# reshape image data for CNN input
train_images = train_images.reshape(train_images.shape[0], 32, 32, 3)
test_images = test_images.reshape(test_images.shape[0], 32, 32, 3)

# one-hot encode labels
train_labels = keras.utils.to_categorical(train_labels, 10)
test_labels = keras.utils.to_categorical(test_labels, 10)

# verify shapes
train_images.shape, test_images.shape

((50000, 32, 32, 3), (10000, 32, 32, 3))
```

```
# Build the CNN model
model = Sequential([
    # First convolutional layer
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3),
    layers.MaxPooling2D((2, 2)),

    # Second convolutional layer
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    # Third convolutional layer
    layers.Conv2D(128, (3, 3), activation='relu'),

    # Flatten the 3D output to 1D
    layers.Flatten(),

    # Dense (fully connected) layer
    layers.Dense(128, activation='relu'),

    # Output layer
    layers.Dense(10, activation='softmax')
])

/usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional.py:11: DeprecationWarning: The class 'Conv2D' is deprecated. Use 'keras.layers.Conv2D' instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=10, batch_size=32,
                     validation_data=(test_images, test_labels))
```

```
Epoch 1/10
782/782  72s 89ms/step - accuracy: 0.3639 - loss: 0.8571
Epoch 2/10
782/782  68s 87ms/step - accuracy: 0.5844 - loss: 0.8092
Epoch 3/10
782/782  85s 90ms/step - accuracy: 0.6675 - loss: 0.8310
Epoch 4/10
782/782  67s 86ms/step - accuracy: 0.7030 - loss: 0.8519
Epoch 5/10
782/782  84s 88ms/step - accuracy: 0.7438 - loss: 0.8571
Epoch 6/10
782/782  69s 89ms/step - accuracy: 0.7614 - loss: 0.8571
Epoch 7/10
782/782  69s 88ms/step - accuracy: 0.7884 - loss: 0.8571
Epoch 8/10
782/782  68s 86ms/step - accuracy: 0.8092 - loss: 0.8571
Epoch 9/10
782/782  83s 88ms/step - accuracy: 0.8310 - loss: 0.8571
Epoch 10/10
782/782  67s 86ms/step - accuracy: 0.8519 - loss: 0.8571
```

```
#Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=0)
print('\nTest accuracy:', test_acc)
```



```
313/313 - 4s - 12ms/step - accuracy: 0.7213 - loss: 0.8571
```


```
Test accuracy: 0.7213000059127808
```

```
# plot training & validation accuracy values
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Test Accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```

```
# plot training & validation loss values
plt.plot(history.history['loss'], label='Train Loss')
```

```
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```





```
# make predictions
predictions = model.predict(test_images)

#display the first 5 predictions
for i in range(5):
    plt.imshow(test_images[i].reshape(32, 32, 3), cmap='gray')
    plt.title(f'Predicted: {predictions[i].argmax()}, Actual: {y_te
    plt.show()
```

