

# Logistic-SGD

December 28, 2025

## 1 Mini Project: Predicting Malaria Diagnosis using SGD

This project involves building a machine learning model to predict whether a person is diagnosed with malaria. Logistic Regression, optimized with Stochastic Gradient Descent (SGD), will be implemented to solve this binary classification problem. We will apply Stochastic Gradient Descent (SGD) to predict the likelihood of a person being diagnosed with malaria based on a set of features in a given dataset. We will implement Logistic Regression and optimize the parameters using SGD. The dataset may include features like symptoms, environmental factors, or measurements that correlate with malaria infection. The task is a binary classification problem (Malaria = 1, No Malaria = 0).

```
[2]: #import necessary libraries
import pandas as pd
from sklearn.preprocessing import StandardScaler

# load the dataset
# https://www.kaggle.com/datasets/programmer3/malaria-diagnosis-dataset
df = pd.read_csv('Malaria_Dataset.csv')
# display the first few rows of the dataset
print(df.head())
```

	IP_Number	Age	Sex	Residence_Area	DOA	Discharge_Date	\
0	14xxxx31	52	Female	Mangalore	31-10-2015 20:42	05-11-2015 05:16	
1	28xxxx34	75	Female	Shimoga	03-02-2015 23:28	13-02-2015 19:27	
2	96xxxx43	30	Female	Mangalore	15-11-2019 12:31	19-11-2019 14:31	
3	49xxxx87	89	Female	Mangalore	17-05-2017 17:50	23-05-2017 13:22	
4	48xxxx10	62	Male	Shimoga	26-06-2015 15:29	27-06-2015 23:35	

	Fever	Headache	Abdominal_Pain	General_Body_Malaise	...	Vomiting	\
0	0	0	0		1	...	0
1	1	0	1		1	...	0
2	1	1	1		1	...	0
3	0	0	0		0	...	1
4	0	1	0		1	...	1

	Confusion	Backache	Chest_Pain	Coughing	Joint_Pain	Primary_Code	\
0	0	1	0	0	0	B50.9	
1	1	0	1	1	1	B50.9	

```

2      1      1      1      0      1      B50.9
3      1      1      1      0      1      B54
4      1      0      0      0      0      B51.0

```

			Diagnosis_Type	Target	Risk_Score
0		Mixed Malaria	Infection	0	3
1		Mixed Malaria	Infection	1	11
2			Mixed Malaria Infection	1	13
3	Plasmodium vivax	Malaria without complication		0	5
4	Plasmodium falciparum	Malaria without complica...		1	8

[5 rows x 21 columns]

[3]: *#summary of variables*  
`print(df.info())`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1622 entries, 0 to 1621
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   IP_Number        1622 non-null   object 
 1   Age              1622 non-null   int64  
 2   Sex              1622 non-null   object 
 3   Residence_Area  1622 non-null   object 
 4   DOA              1622 non-null   object 
 5   Discharge_Date  1622 non-null   object 
 6   Fever             1622 non-null   int64  
 7   Headache          1622 non-null   int64  
 8   Abdominal_Pain   1622 non-null   int64  
 9   General_Body_Malaise 1622 non-null   int64  
 10  Dizziness         1622 non-null   int64  
 11  Vomiting          1622 non-null   int64  
 12  Confusion          1622 non-null   int64  
 13  Backache           1622 non-null   int64  
 14  Chest_Pain         1622 non-null   int64  
 15  Coughing           1622 non-null   int64  
 16  Joint_Pain         1622 non-null   int64  
 17  Primary_Code       1622 non-null   object 
 18  Diagnosis_Type    1622 non-null   object 
 19  Target             1622 non-null   int64  
 20  Risk_Score          1622 non-null   int64  
dtypes: int64(14), object(7)
memory usage: 266.2+ KB
None

```

[4]: *# check for missing values and duplicates and display on same table*  
`missing_values = df.isnull().sum()`

```

duplicate_values = df.duplicated().sum()

summary_df = pd.DataFrame({'Missing Values': missing_values, 'Duplicate Values':
    ↪ [duplicate_values]*len(missing_values)})

print(summary_df)

```

	Missing Values	Duplicate Values
IP_Number	0	0
Age	0	0
Sex	0	0
Residence_Area	0	0
DOA	0	0
Discharge_Date	0	0
Fever	0	0
Headache	0	0
Abdominal_Pain	0	0
General_Body_Malaise	0	0
Dizziness	0	0
Vomiting	0	0
Confusion	0	0
Backache	0	0
Chest_Pain	0	0
Coughing	0	0
Joint_Pain	0	0
Primary_Code	0	0
Diagnosis_Type	0	0
Target	0	0
Risk_Score	0	0

[5]: # summary statistics of the dataset  
print(df.describe())

	Age	Fever	Headache	Abdominal_Pain	\
count	1622.000000	1622.000000	1622.000000	1622.000000	
mean	44.705302	0.474723	0.507398	0.496917	
std	26.025626	0.499515	0.500099	0.500145	
min	1.000000	0.000000	0.000000	0.000000	
25%	22.000000	0.000000	0.000000	0.000000	
50%	44.000000	0.000000	1.000000	0.000000	
75%	68.000000	1.000000	1.000000	1.000000	
max	89.000000	1.000000	1.000000	1.000000	

	General_Body_Malaise	Dizziness	Vomiting	Confusion	\
count	1622.000000	1622.000000	1622.000000	1622.000000	
mean	0.479038	0.504316	0.500000	0.479655	
std	0.499714	0.500136	0.500154	0.499740	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	

```

50%          0.000000    1.000000    0.500000    0.000000
75%          1.000000    1.000000    1.000000    1.000000
max          1.000000    1.000000    1.000000    1.000000

      Backache   Chest_Pain   Coughing   Joint_Pain   Target \
count  1622.000000  1622.000000  1622.000000  1622.000000  1622.000000
mean    0.509864    0.500617    0.483970    0.511714    0.719482
std     0.500057    0.500154    0.499897    0.500017    0.449391
min     0.000000    0.000000    0.000000    0.000000    0.000000
25%    0.000000    0.000000    0.000000    0.000000    0.000000
50%    1.000000    1.000000    0.000000    1.000000    1.000000
75%    1.000000    1.000000    1.000000    1.000000    1.000000
max    1.000000    1.000000    1.000000    1.000000    1.000000

      Risk_Score
count  1622.000000
mean    8.405055
std     3.088743
min     0.000000
25%    6.000000
50%    9.000000
75%   11.000000
max   17.000000

```

```
[6]: # convert sex to numeric
# create a mapping dictionary
gender_mapping = {'Female': 0, 'Male': 1}
# apply the mapping to the 'Sex' column
df['Sex'] = df['Sex'].map(gender_mapping)

# show df
print(df.head())

```

	IP_Number	Age	Sex	Residence_Area	DOA	Discharge_Date	
0	14xxxx31	52	0	Mangalore	31-10-2015 20:42	05-11-2015 05:16	
1	28xxxx34	75	0	Shimoga	03-02-2015 23:28	13-02-2015 19:27	
2	96xxxx43	30	0	Mangalore	15-11-2019 12:31	19-11-2019 14:31	
3	49xxxx87	89	0	Mangalore	17-05-2017 17:50	23-05-2017 13:22	
4	48xxxx10	62	1	Shimoga	26-06-2015 15:29	27-06-2015 23:35	

	Fever	Headache	Abdominal_Pain	General_Body_Malaise	...	Vomiting	
0	0	0	0		1	...	0
1	1	0	1		1	...	0
2	1	1	1		1	...	0
3	0	0	0		0	...	1
4	0	1	0		1	...	1

```
Confusion  Backache   Chest_Pain   Coughing   Joint_Pain   Primary_Code \

```

```

0      0      1      0      0      0      B50.9
1      1      0      1      1      1      B50.9
2      1      1      1      0      1      B50.9
3      1      1      1      0      1      B54
4      1      0      0      0      0      B51.0

```

	Diagnosis_Type	Target	Risk_Score
0	Mixed Malaria Infection	0	3
1	Mixed Malaria Infection	1	11
2	Mixed Malaria Infection	1	13
3	Plasmodium vivax Malaria without complication	0	5
4	Plasmodium falciparum Malaria without complica...	1	8

[5 rows x 21 columns]

```
[7]: # drop unnecessary columns [IP_Number, Residence_Area, DOA, Discharge_Date, ↴
    ↴Primary_Code, Diagnosis_Type]
df = df.drop(['IP_Number', 'Residence_Area', 'DOA', 'Discharge_Date', ↴
    ↴'Primary_Code', 'Diagnosis_Type'], axis=1)
print(df.head())
```

	Age	Sex	Fever	Headache	Abdominal_Pain	General_Body_Malaise	Dizziness	\
0	52	0	0	0	0		1	0
1	75	0	1	0	1		1	1
2	30	0	1	1	1		1	0
3	89	0	0	0	0		0	1
4	62	1	0	1	0		1	0

	Vomiting	Confusion	Backache	Chest_Pain	Coughing	Joint_Pain	Target	\
0	0	0	1	0	0	0	0	0
1	0	1	0	1	1	1	1	1
2	0	1	1	1	0	1	1	1
3	1	1	1	1	0	1	0	0
4	1	1	0	0	0	0	0	1

	Risk_Score
0	3
1	11
2	13
3	5
4	8

```
[8]: # Prepare features and target variable
X = df.drop('Target', axis=1) # Features
y = df['Target'] # Target variable
```

```
[9]: # Standardizing the dataset since SGD is sensitive to feature scaling
from sklearn.preprocessing import StandardScaler
```

```

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

[10]: # split the dataset into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, u
↪random_state=42)

[11]: # implement logistic regression
import numpy as np

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

class SgdLogisticRegression:
    def __init__(self, lr=0.01, epochs=100, batch_size=32, verbose=False):
        self.lr = lr
        self.epochs = epochs
        self.batch_size = batch_size
        self.verbose = verbose
        self.loss_history = []

    def fit(self, X, y):
        X = np.asarray(X, dtype=float)
        y = np.asarray(y, dtype=float)
        n_samples, n_features = X.shape
        rng = np.random.default_rng(42)
        self.coef_ = np.zeros(n_features)
        self.intercept_ = 0.0
        self.loss_history = []

        for epoch in range(1, self.epochs + 1):
            idx = rng.permutation(n_samples)
            X_shuffled = X[idx]
            y_shuffled = y[idx]

            for start in range(0, n_samples, self.batch_size):
                end = start + self.batch_size
                X_batch = X_shuffled[start:end]
                y_batch = y_shuffled[start:end]

                preds = sigmoid(X_batch.dot(self.coef_) + self.intercept_)
                error = preds - y_batch
                grad_w = X_batch.T.dot(error) / len(y_batch)
                grad_b = error.mean()

                self.coef_ -= self.lr * grad_w

```

```

        self.intercept_ -= self.lr * grad_b

        # compute log-loss for monitoring
        preds_full = sigmoid(X.dot(self.coef_) + self.intercept_)
        eps = 1e-15
        loss = -np.mean(y * np.log(preds_full + eps) + (1 - y) * np.log(1 - preds_full + eps))
        self.loss_history.append(loss)
        if self.verbose and epoch % 10 == 0:
            print(f"Epoch {epoch}/{self.epochs} - Log-loss: {loss:.6f}")

    def predict_proba(self, X):
        X = np.asarray(X, dtype=float)
        z = X.dot(self.coef_) + self.intercept_
        probs = sigmoid(z)
        return np.vstack([1 - probs, probs]).T

    def predict(self, X, threshold=0.5):
        probs = self.predict_proba(X)[:, 1]
        return (probs >= threshold).astype(int)

# training configuration
epochs = 100
learning_rate = 0.01
batch_size = 32

# instantiate and train the model
sgd_model = SgdLogisticRegression(lr=learning_rate, epochs=epochs,
                                   batch_size=batch_size, verbose=True)
sgd_model.fit(X_train, y_train)
loss_history = sgd_model.loss_history

```

Epoch 10/100 - Log-loss: 0.326674  
 Epoch 20/100 - Log-loss: 0.251487  
 Epoch 30/100 - Log-loss: 0.215209  
 Epoch 40/100 - Log-loss: 0.192835  
 Epoch 50/100 - Log-loss: 0.177276  
 Epoch 60/100 - Log-loss: 0.165648  
 Epoch 70/100 - Log-loss: 0.156489  
 Epoch 80/100 - Log-loss: 0.149040  
 Epoch 90/100 - Log-loss: 0.142803  
 Epoch 100/100 - Log-loss: 0.137468

[12]:

```
# log loss evolution
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
```


```



plt.plot(range(1, epochs + 1), loss_history, marker='o')
plt.title('Log-Loss Evolution over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Log-Loss')
plt.grid()
plt.show()

# confusion matrix and classification report
# make predictions
y_pred = sgd_model.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Malaria', 'Malaria'], yticklabels=['No Malaria', 'Malaria'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion Matrix')
plt.show()

# actual vs predicted
plt.figure(figsize=(10, 6))
plt.scatter(range(len(y_test)), y_test, color='blue', label='Actual', alpha=0.6, s=5)
plt.scatter(range(len(y_pred)), y_pred, color='red', label='Predicted', alpha=0.6, s=5)
plt.title('Actual vs Predicted Values')
plt.legend()
plt.show()

```





```
[13]: # make predictions
y_pred = sgd_model.predict(X_test)

# evaluate the model
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.4f}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9938  
Classification Report:

	precision	recall	f1-score	support
0	1.00	0.98	0.99	92
1	0.99	1.00	1.00	233
accuracy			0.99	325
macro avg	1.00	0.99	0.99	325
weighted avg	0.99	0.99	0.99	325

## 1.1 Report

This report summarizes the development and evaluation of a logistic regression model designed to predict malaria diagnosis using data from the Malaria\_Dataset.csv file. The dataset was preprocessed by removing non-predictive administrative columns (IP\_Number, Residence\_Area, DOA, Discharge\_Date, Primary\_Code, and Diagnosis\_Type), encoding the categorical Sex feature as 0 for Female and 1 for Male, and standardizing all numerical features with StandardScaler. An 80/20 train-test split was applied with a fixed random state of 42 for reproducibility, resulting in a test set of 325 samples.

A logistic regression model was implemented from scratch, utilizing the sigmoid activation function and optimized via mini-batch stochastic gradient descent on the binary cross-entropy loss function. Training was conducted over 100 epochs with a learning rate of 0.01 and a batch size of 32, during which the log-loss steadily declined from approximately 0.60 to 0.14, demonstrating effective convergence.

The model achieved outstanding performance with an accuracy of 0.9938. The confusion matrix revealed 90 true negatives, 233 true positives, 2 false positives, and 0 false negatives, indicating near-perfect detection of malaria cases (recall of 1.00 for the positive class) and very high precision overall. The classification report further confirmed precision, recall, and F1-scores near 1.00 for both classes, with macro and weighted averages around 0.99.

Overall, the SGD-trained logistic regression model demonstrated exceptional predictive capability on this dataset split, making it highly reliable for identifying malaria cases with minimal errors.