

Time Series Analysis of Daily Total Female Births Dataset From Kaggle

Bett

2025-12-21

```
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo

library(tseries)
```

1. Load the data Make sure you've downloaded 'DailyTotalFemaleBirths.csv' from Kaggle and saved in working directory.

```
data <- read.csv("DailyTotalFemaleBirths.csv")
```

2. Inspect the data structure

```
str(data)

## 'data.frame': 365 obs. of 2 variables:
## $ Date : chr "1959-01-01" "1959-01-02" "1959-01-03" "1959-01-04" ...
## $ Births: int 35 32 30 31 44 29 45 43 38 27 ...
```

```
head(data)
```

```
##          Date Births
## 1 1959-01-01     35
## 2 1959-01-02     32
## 3 1959-01-03     30
## 4 1959-01-04     31
## 5 1959-01-05     44
## 6 1959-01-06     29
```

The dataset typically has columns: 'Date' and 'Births'. Convert the 'Date' column to Date class

```
data$Date <- as.Date(data$Date, format="%Y-%m-%d")
```

Create a time series object Since these are daily births for 1959, frequency = 365 is often used for daily data (non-leap year)

```
births_ts <- ts(data$Births, start=c(1959,1), frequency=365)
```

2. Test for stationarity We can use the Augmented Dickey-Fuller (ADF) test:

```
adf_result <- adf.test(births_ts)
```

```
## Warning in adf.test(births_ts): p-value smaller than printed p-value
```

```

adf_result

##
##  Augmented Dickey-Fuller Test
##
## data:  births_ts
## Dickey-Fuller = -5.1042, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary

```

Interpretation:

If the p-value is small (<0.05), we can reject the null hypothesis of non-stationarity. In this dataset, the daily female births data is often already stationary or nearly so. If not stationary, we would consider differencing:
`births_ts_diff <- diff(births_ts)`

(Check stationarity on differenced series if needed) `adf.test(births_ts_diff)`

3. Split the data into training and testing sets for model evaluation Let's hold out the last 30 days as test data.

```

train_length <- length(births_ts) - 30
train_ts <- window(births_ts, end=c(1959,(train_length/365)*365))
test_ts <- window(births_ts, start=c(1959, ((train_length+1)/365)*365))

```

4. Fit an ARIMA model using `auto.arima` on the training set

```

fit <- auto.arima(train_ts)
summary(fit)

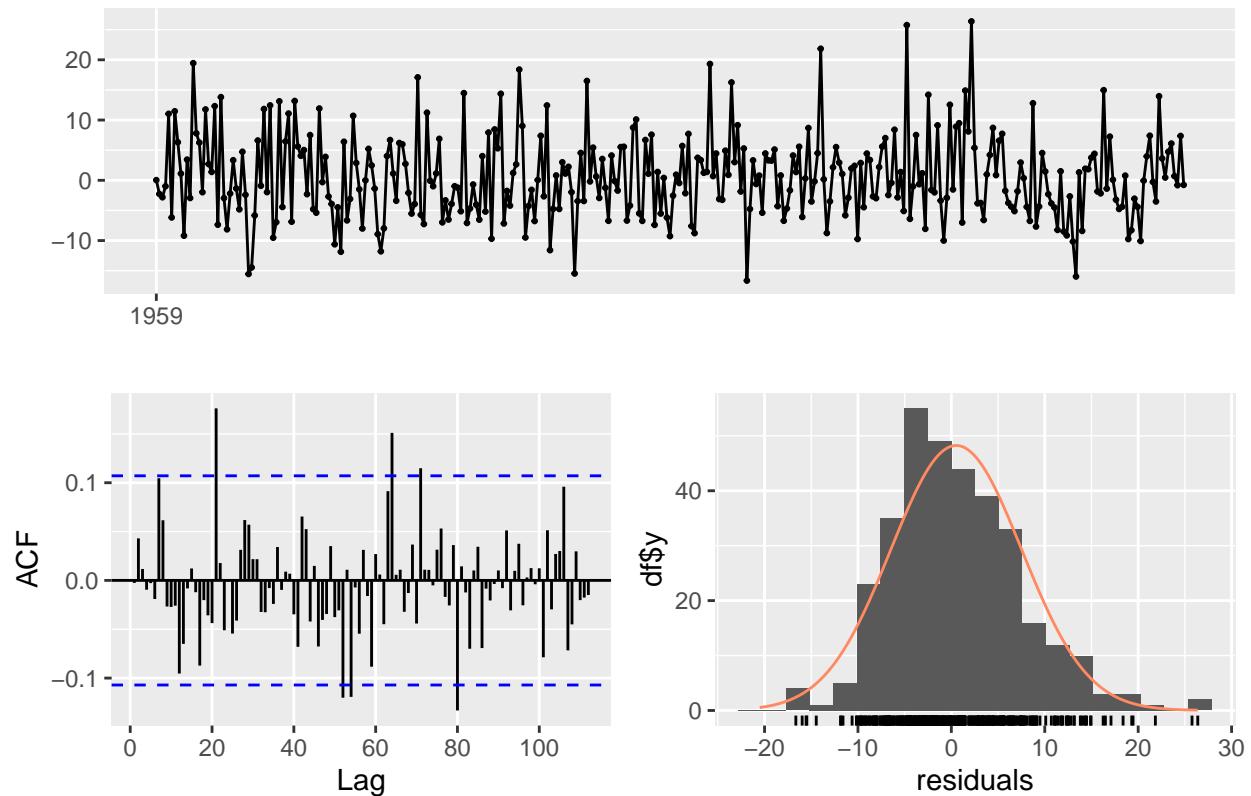
## Series: train_ts
## ARIMA(0,1,2)
##
## Coefficients:
##             ma1      ma2
##           -0.8573  -0.0964
## s.e.    0.0523   0.0525
##
## sigma^2 = 49.73:  log likelihood = -1126.5
## AIC=2259  AICc=2259.08  BIC=2270.44
##
## Training set error measures:
##          ME      RMSE      MAE      MPE      MAPE MASE        ACF1
## Training set 0.5224927 7.020607 5.465459 -1.449279 13.26883  NaN -0.002565009

```

Check the residuals of the fitted model

```
checkresiduals(fit)
```

Residuals from ARIMA(0,1,2)



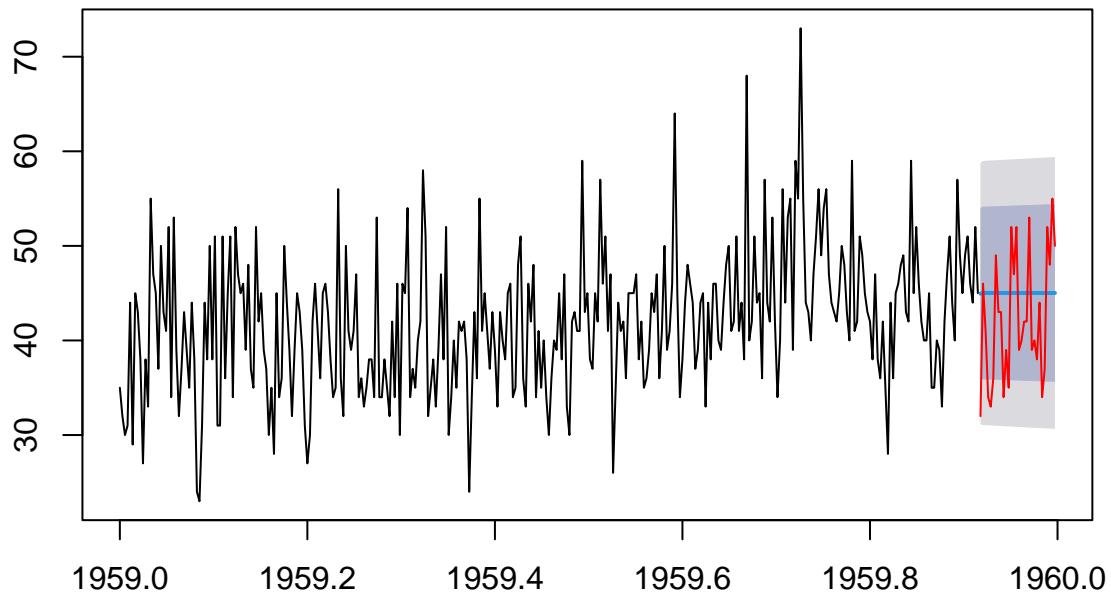
```
##  
## Ljung-Box test  
##  
## data: Residuals from ARIMA(0,1,2)  
## Q* = 75.269, df = 65, p-value = 0.1801  
##  
## Model df: 2. Total lags used: 67
```

If residuals look like white noise (no patterns, no autocorrelations), the model is adequate.

5. Forecast on the test period Forecast for the length of the test set (30 days)

```
fcast <- forecast(fit, h=30)
plot(fcast)
lines(test_ts, col="red", type="l") # Add the actual test data in red
```

Forecasts from ARIMA(0,1,2)



Evaluate forecasting accuracy

```
accuracy_metrics <- accuracy(fcast, test_ts)
accuracy_metrics
```

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
## Training set	0.5224927	7.020607	5.465459	-1.449279	13.26883	NaN	-0.002565009
## Test set	-2.7173817	7.160799	6.304155	-9.038007	15.97857	NaN	0.216342192
## Theil's U							
## Training set		NA					
## Test set		0.8474797					

Look at RMSE, MAE, MAPE, etc. to assess forecast accuracy.

6. If satisfied with the model, you can refit using the entire dataset and forecast future values

```
final_fit <- auto.arima(births_ts)
fcast_future <- forecast(final_fit, h=30) # forecasting next 30 days beyond 1959
plot(fcast_future)
```

Forecasts from ARIMA(0,1,2)

